

# Applied Machine Learning

Deep Neural Networks and Backpropagation

# Deep Neural Networks and Backpropagation

- Backpropagation Algorithm
- Training Deep Neural Networks

# Backpropagation Algorithm

- Layer  $i - 1$ : changes in Loss at output with respect to  $\theta^{(i-1)}$

$$\bullet \quad \nabla_{\theta^{(i)}} L = \nabla_{\mathbf{o}^{(D)}} L \times \mathbf{J}_{\mathbf{o}^{(D)}; \mathbf{u}^{(D)}} \times \dots \times \mathbf{J}_{\mathbf{o}^{(i+1)}; \mathbf{u}^{(i+1)}} \times \mathbf{J}_{\mathbf{o}^{(i)}; \theta^{(i)}}$$

$$\mathbf{v}^{(D)} = \nabla_{\mathbf{o}^{(D)}} L$$

$$\nabla_{\theta^{(D)}} L = \mathbf{v}^{(D)} \times \mathbf{J}_{\mathbf{o}^{(D)}; \theta^{(D)}}$$

$$\vdots$$

$$\mathbf{v}^{(i)} = \mathbf{v}^{(i+1)} \times \mathbf{J}_{\mathbf{o}^{(i+1)}; \mathbf{u}^{(i+1)}}$$

$$\bullet \quad \begin{aligned} \nabla_{\theta^{(i)}} L &= \mathbf{v}^{(i)} \times \mathbf{J}_{\mathbf{o}^{(i)}; \theta^{(i)}} \\ &\vdots \end{aligned}$$

- Backpropagation

- Forward pass:

$$\bullet \quad \mathbf{u}^{(1)} = \mathbf{x}$$

- for each layer  $i$  from 1 to D:

$$\bullet \quad \mathbf{u}^{(i+1)} = \mathbf{o}^{(i)}(\mathbf{u}^{(i)}, \theta^{(i)})$$

- Backward pass:

$$\mathbf{v}^{(D)} = \nabla_{\mathbf{o}^{(D)}} L$$

$$\bullet \quad \nabla_{\theta^{(D)}} L = \mathbf{v}^{(D)} \times \mathbf{J}_{\mathbf{o}^{(D)}; \theta^{(D)}}$$

- for each layer  $i$  from D-1 to 1:

$$\mathbf{v}^{(i)} = \mathbf{v}^{(i+1)} \times \mathbf{J}_{\mathbf{o}^{(i+1)}; \mathbf{u}^{(i+1)}}$$

$$\bullet \quad \nabla_{\theta^{(i)}} L = \mathbf{v}^{(i)} \times \mathbf{J}_{\mathbf{o}^{(i)}; \theta^{(i)}}$$

# Parameter Initialization

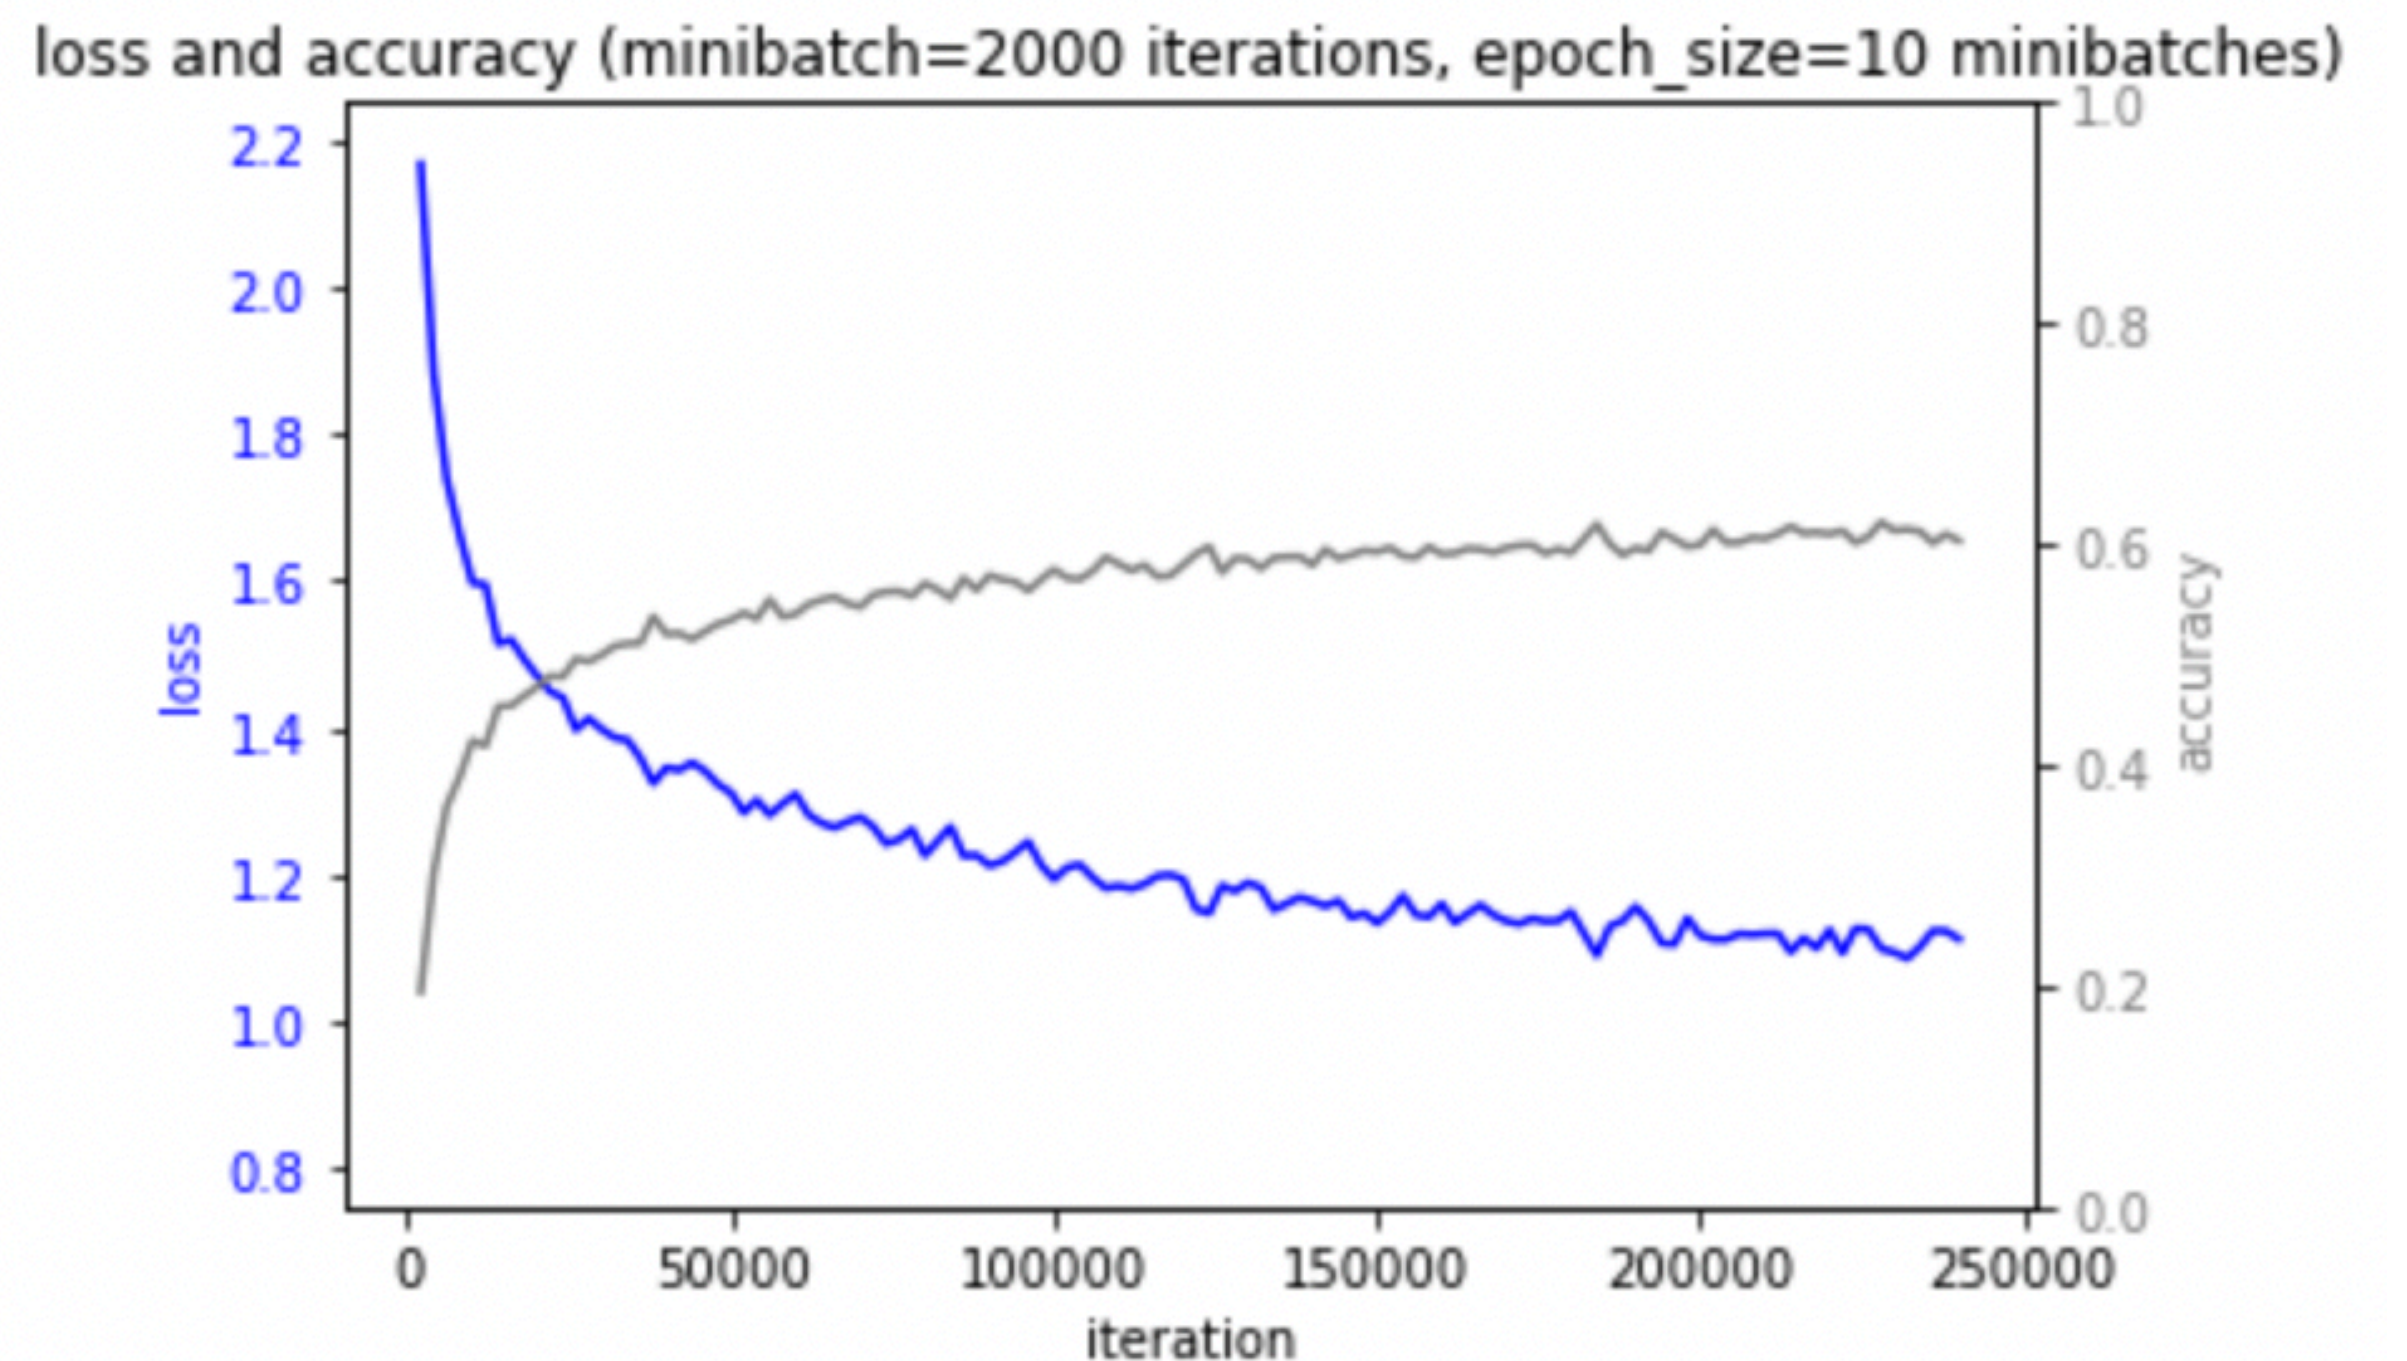
- Initial value of parameters
  - All 0 => bad option
  - Close to their optimal value
  - Each to a random value
    - from zero mean normal distribution with small standard deviation (e.g. 0.01)

# Feature Preprocessing

- The relative distribution of different feature inputs may affect learning
  - Large variations in their standard deviations
    - Weights for inputs with larger standard deviations: larger
      - if stepsize is small, it takes longer to build them
    - Weights for inputs with smaller standard deviations: smaller
      - if stepsize is large, it's unlikely to find them
- Preprocessing:
  - Normalize each input to have zero mean and unit standard deviation
  - Apply domain knowledge

# Training Loss and Accuracy

- Plots: Loss vs Epochs, Accuracy vs Epochs
- Stepsize or Learning Rate
  - smaller learning rate
    - smoother curves => gradual walk towards minimum
    - more likely to fall into non-optimal loss
  - larger learning rate
    - more noisy curves => jumps on the gradient
    - more likely to diverge
- Reducing loss does not necessarily improve accuracy



# Data

- It is hard to successfully train a Deep Neural Network
- Datasets
  - Many times, the more data the better, up to a point

# Units

- Units
  - Redundant units
    - May result in cancellation of some at next layer
  - Dropout
    - before each training step
      - set outputs of some randomly selected units to zero
        - do not consider output units
- Dead units
  - output 0 for every data item => Gradient 0 => dead unit
  - smaller learning rates help
  - many more units



# Gradient Obfuscation

- Gradient Obfuscation
  - poor parameter estimates close to output layer => poor gradient update
  - the more layers above, the worst
- initialize with good estimate
- rescaling
- change connectivity structure

# Training Deep Neural Networks

- Implementation
  - GPUs
  - APIs
    - describe layers, connectivity, gradients, unit functions
    - map onto the available computing resources (e.g., GPU)
    - training, and evaluation

# Deep Neural Networks and Backpropagation

- Backpropagation Algorithm
- Training Deep Neural Networks

# Applied Machine Learning

Deep Neural Networks and Backpropagation