

Applied Machine Learning

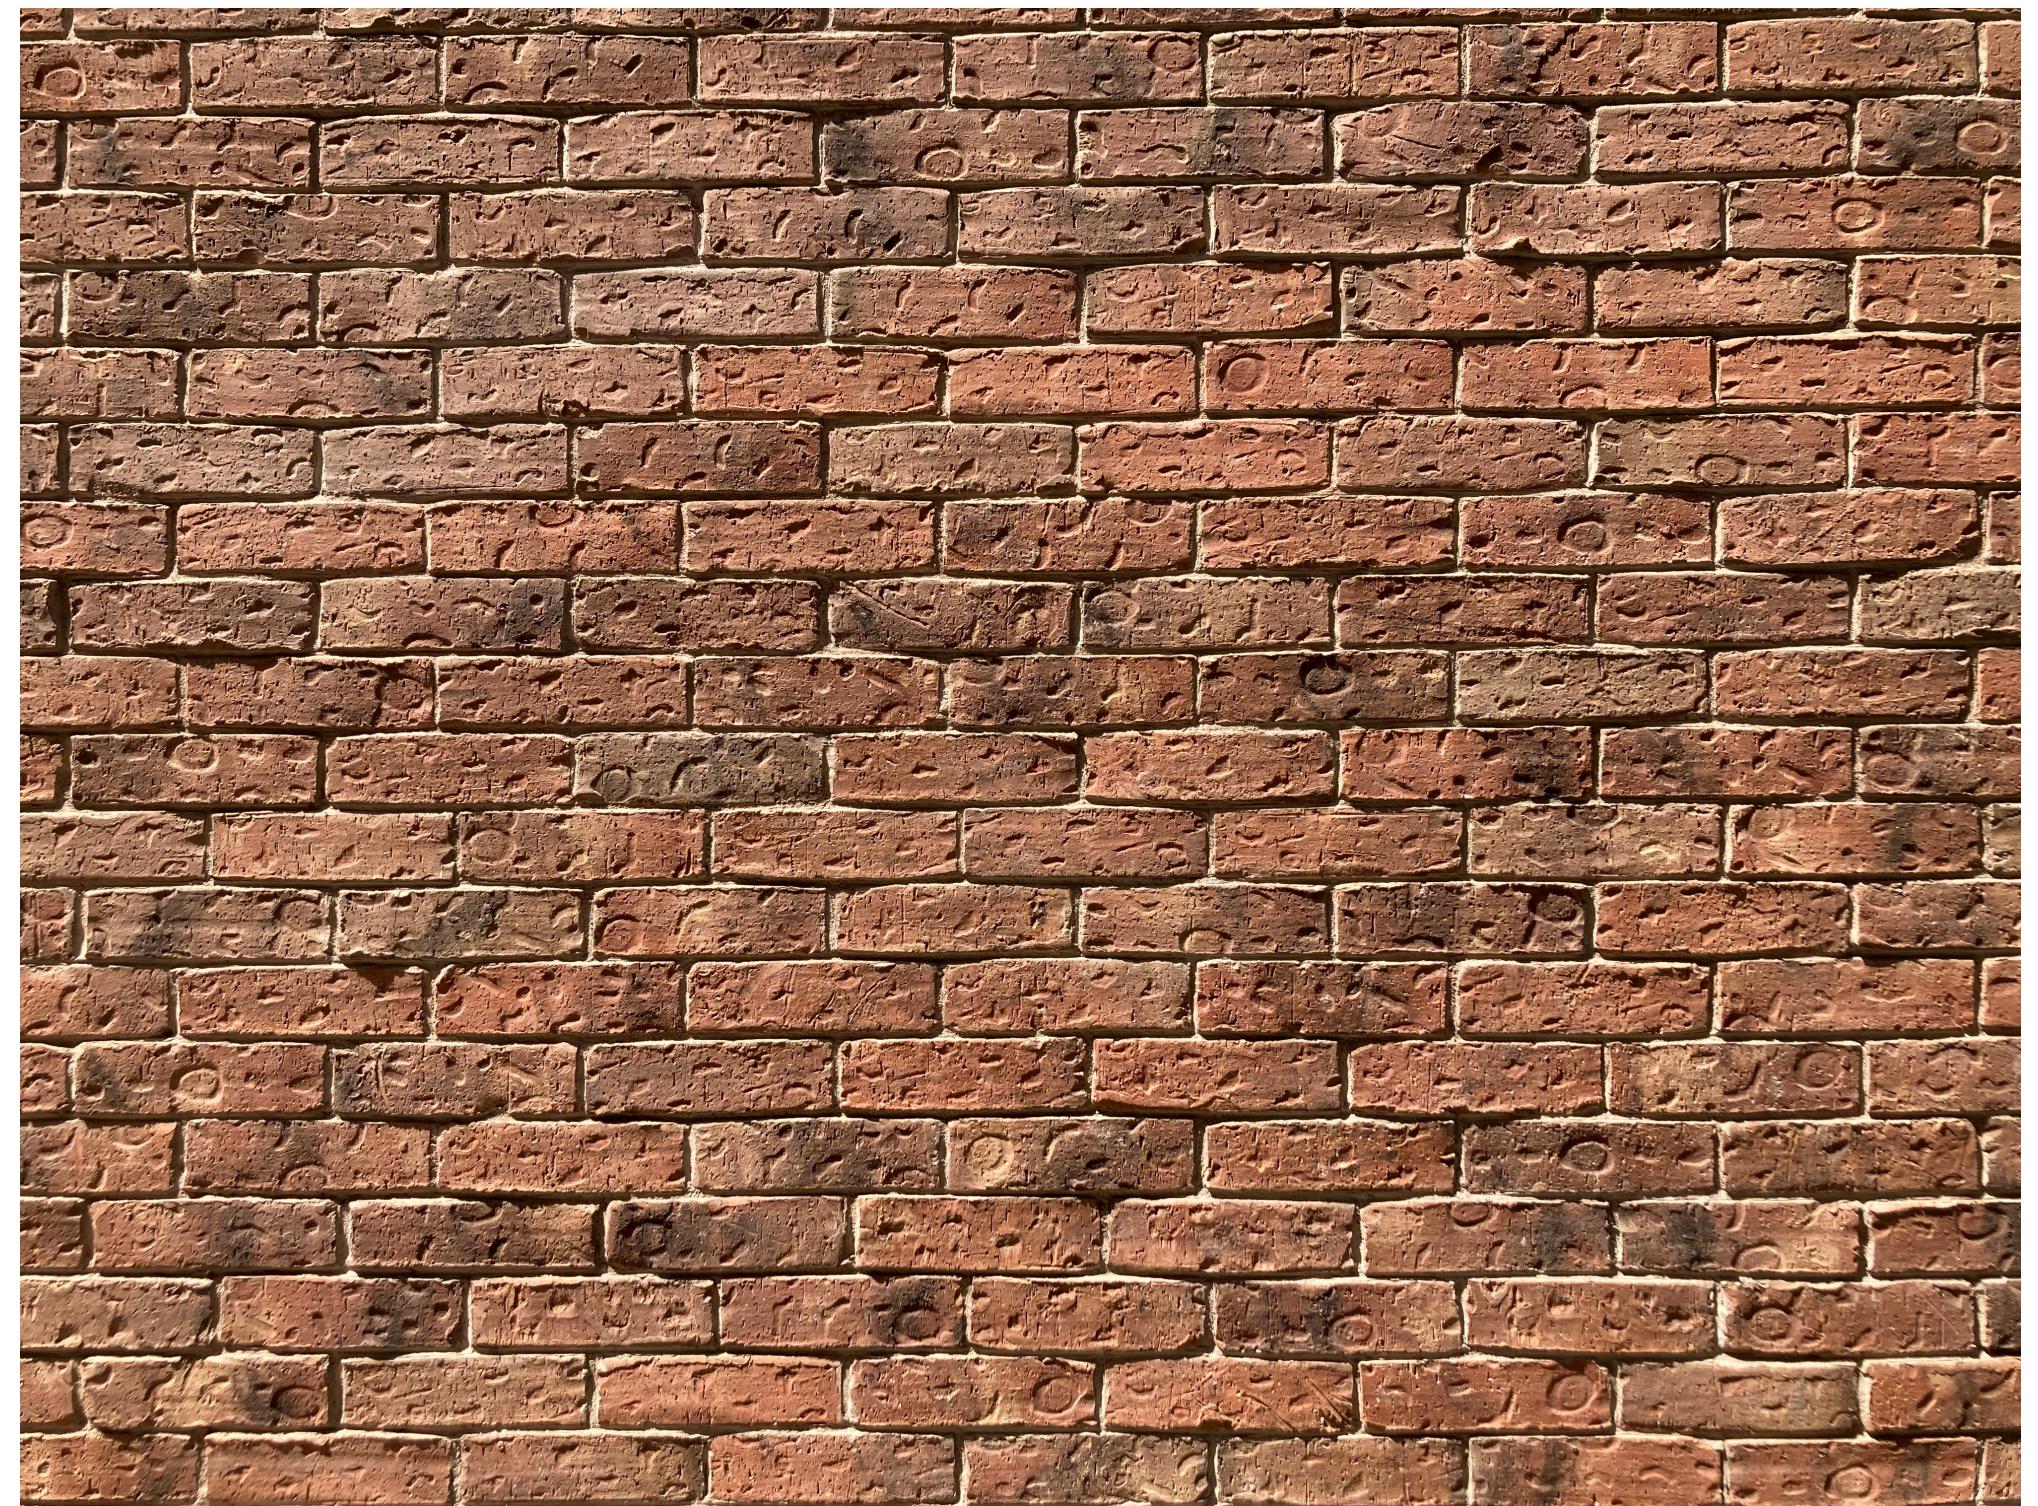
Convolutional Layer

Convolutional Layer

- Feature Maps
- From convolutional kernel to linear unit
- Convolutional Layer
- Pooling Layer

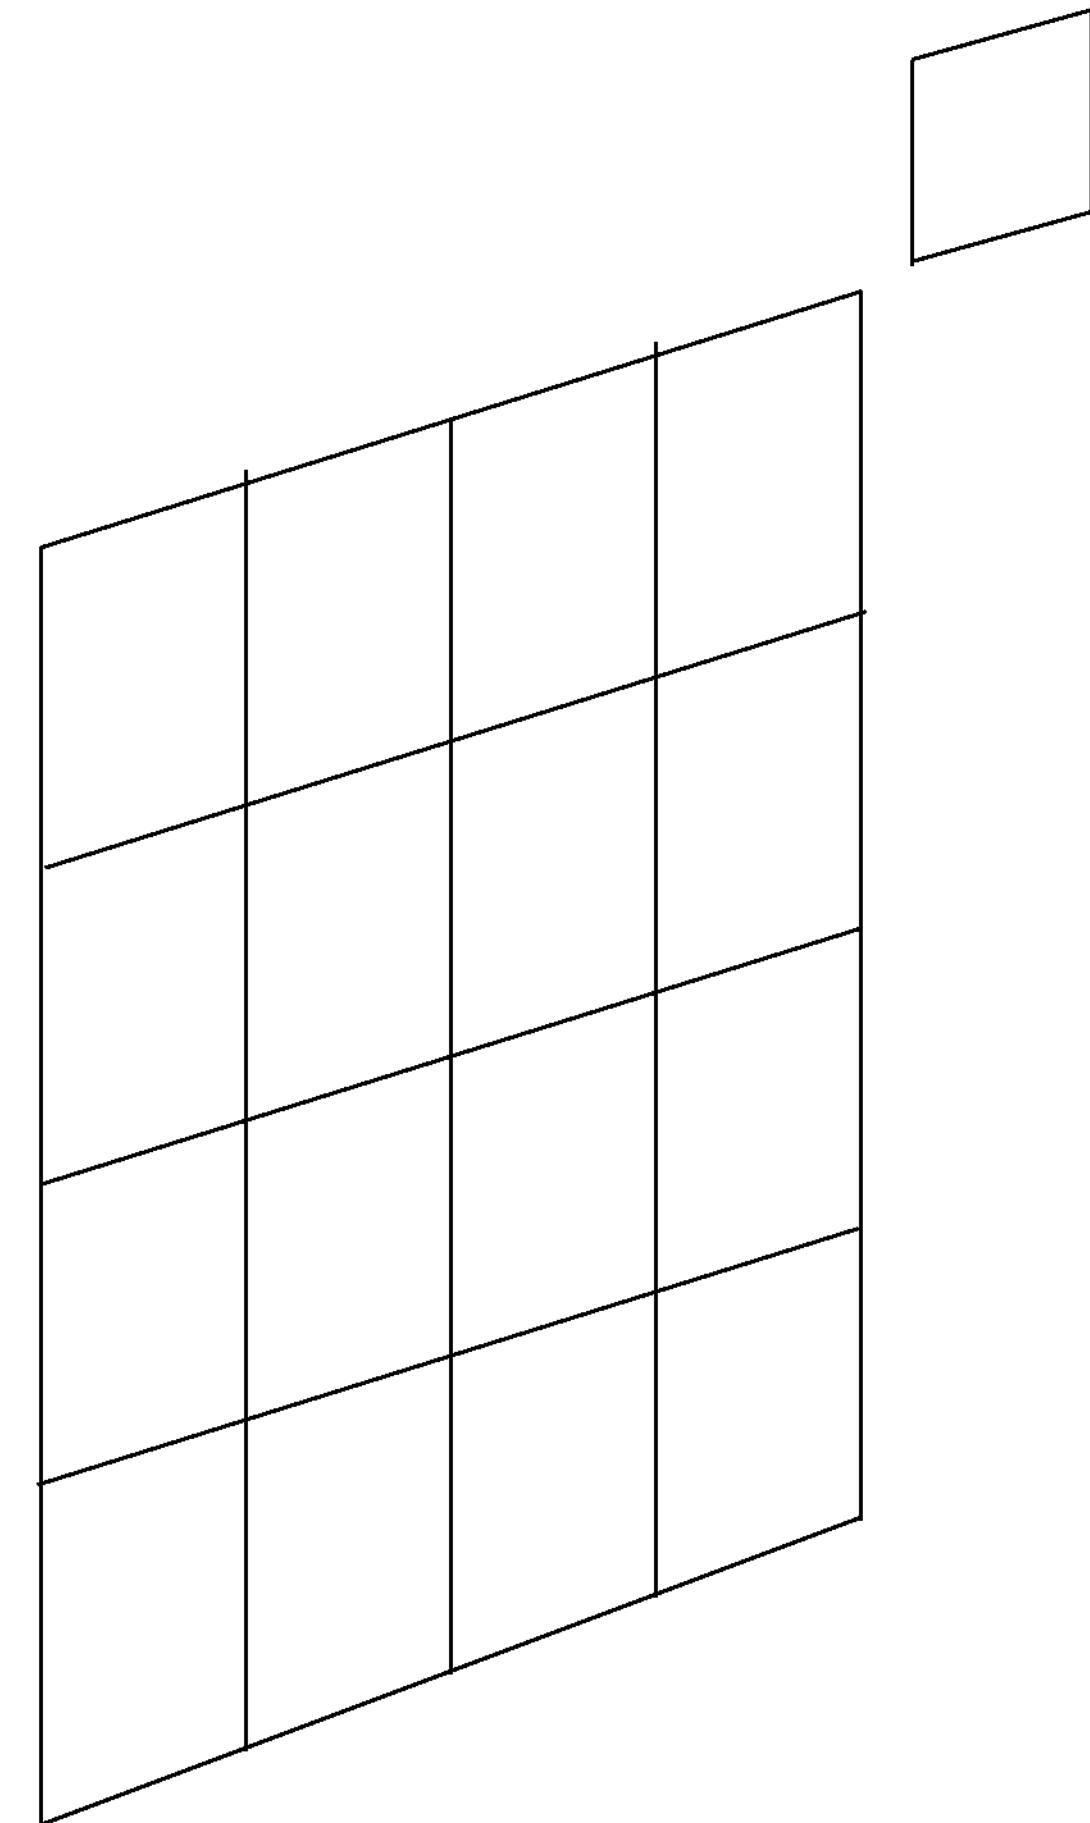
3-Dimensional Convolution

- Color image I
 - 3 planar channels: Red, Green, Blue (RGB): I_r, I_g, I_b
 - stacked in 3D array or Tensor, 1 slice per channel
- 3D convolution
 - Similar to 2D convolution
 - Kernel: 3D array or Tensor: 1 slice per channel
 - $N_{u,v}$: weighted sum within pixels defined by 3D Kernel around pixel (u, v)
 - 2-dimensional
- same interpretation as 2D convolution
- same properties
- similar padding and stride strategies



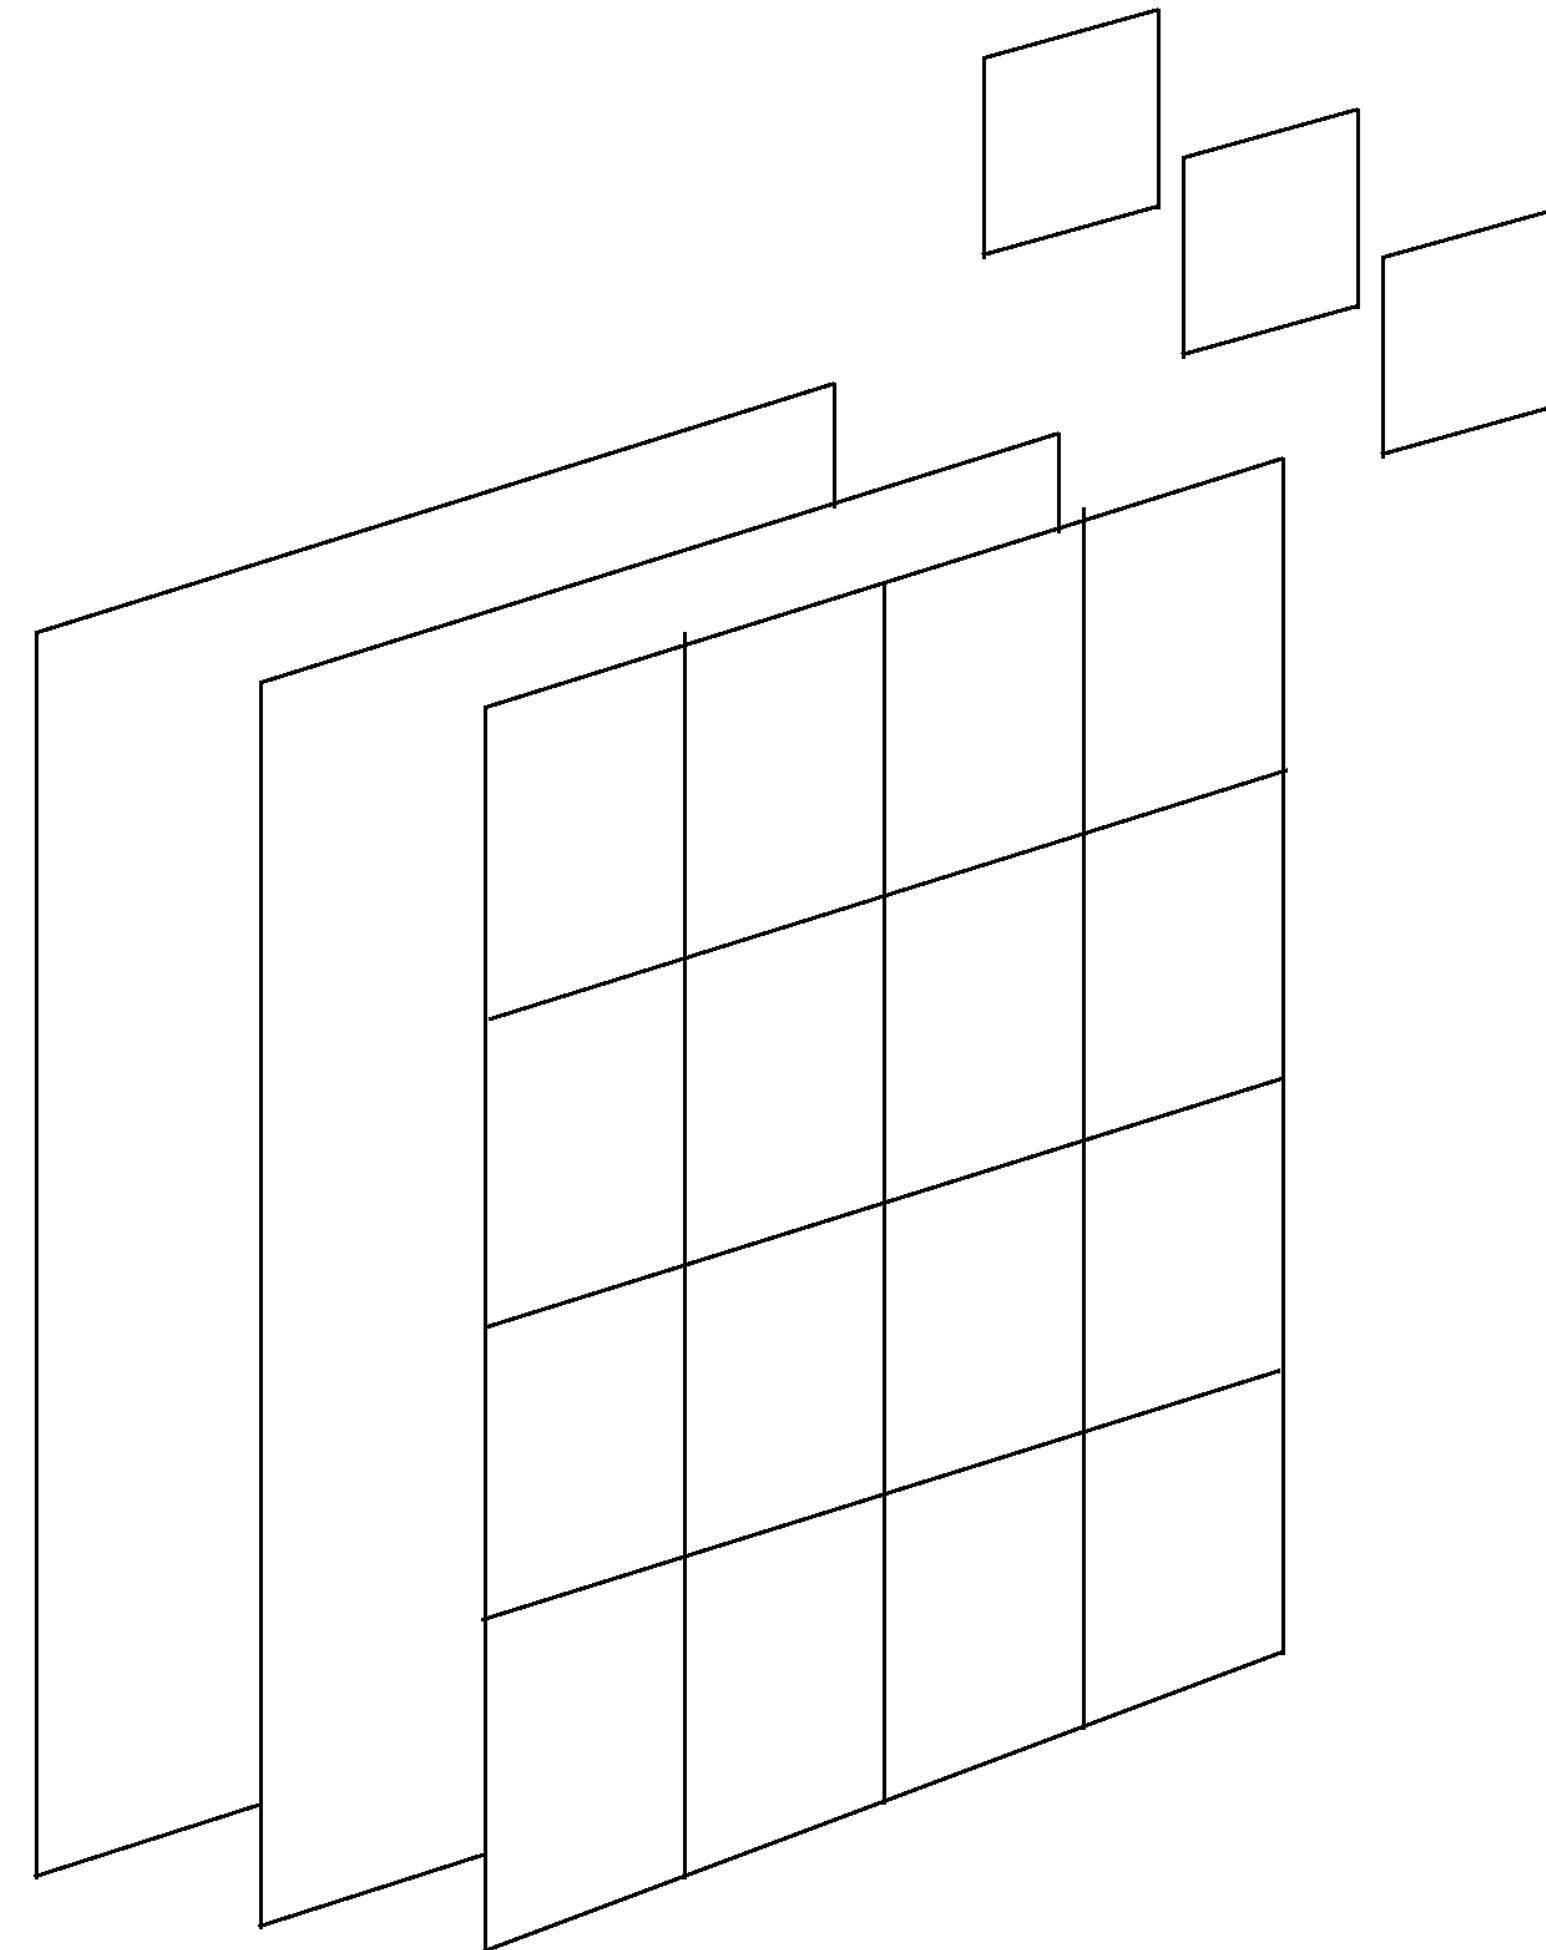
Feature Map

- Input: Image I of size $x \times y$ with D channels
 - 3-D: $x \times y \times D$
- Collection of kernels K : stack of L Kernels
 - K_l : Kernel for pattern $l \in \{1 \dots L\}$
 - dimensions: same as input I : $k_i \times k_j \times D$
 - dimensions: $k_i \times k_j \times D \times L$
- Output: Stack of L responses: $N_l = I \otimes K_l$
 - 3-dimensional: $x_n \times y_n \times L$
 - x_n, y_n : depend on strides and padding



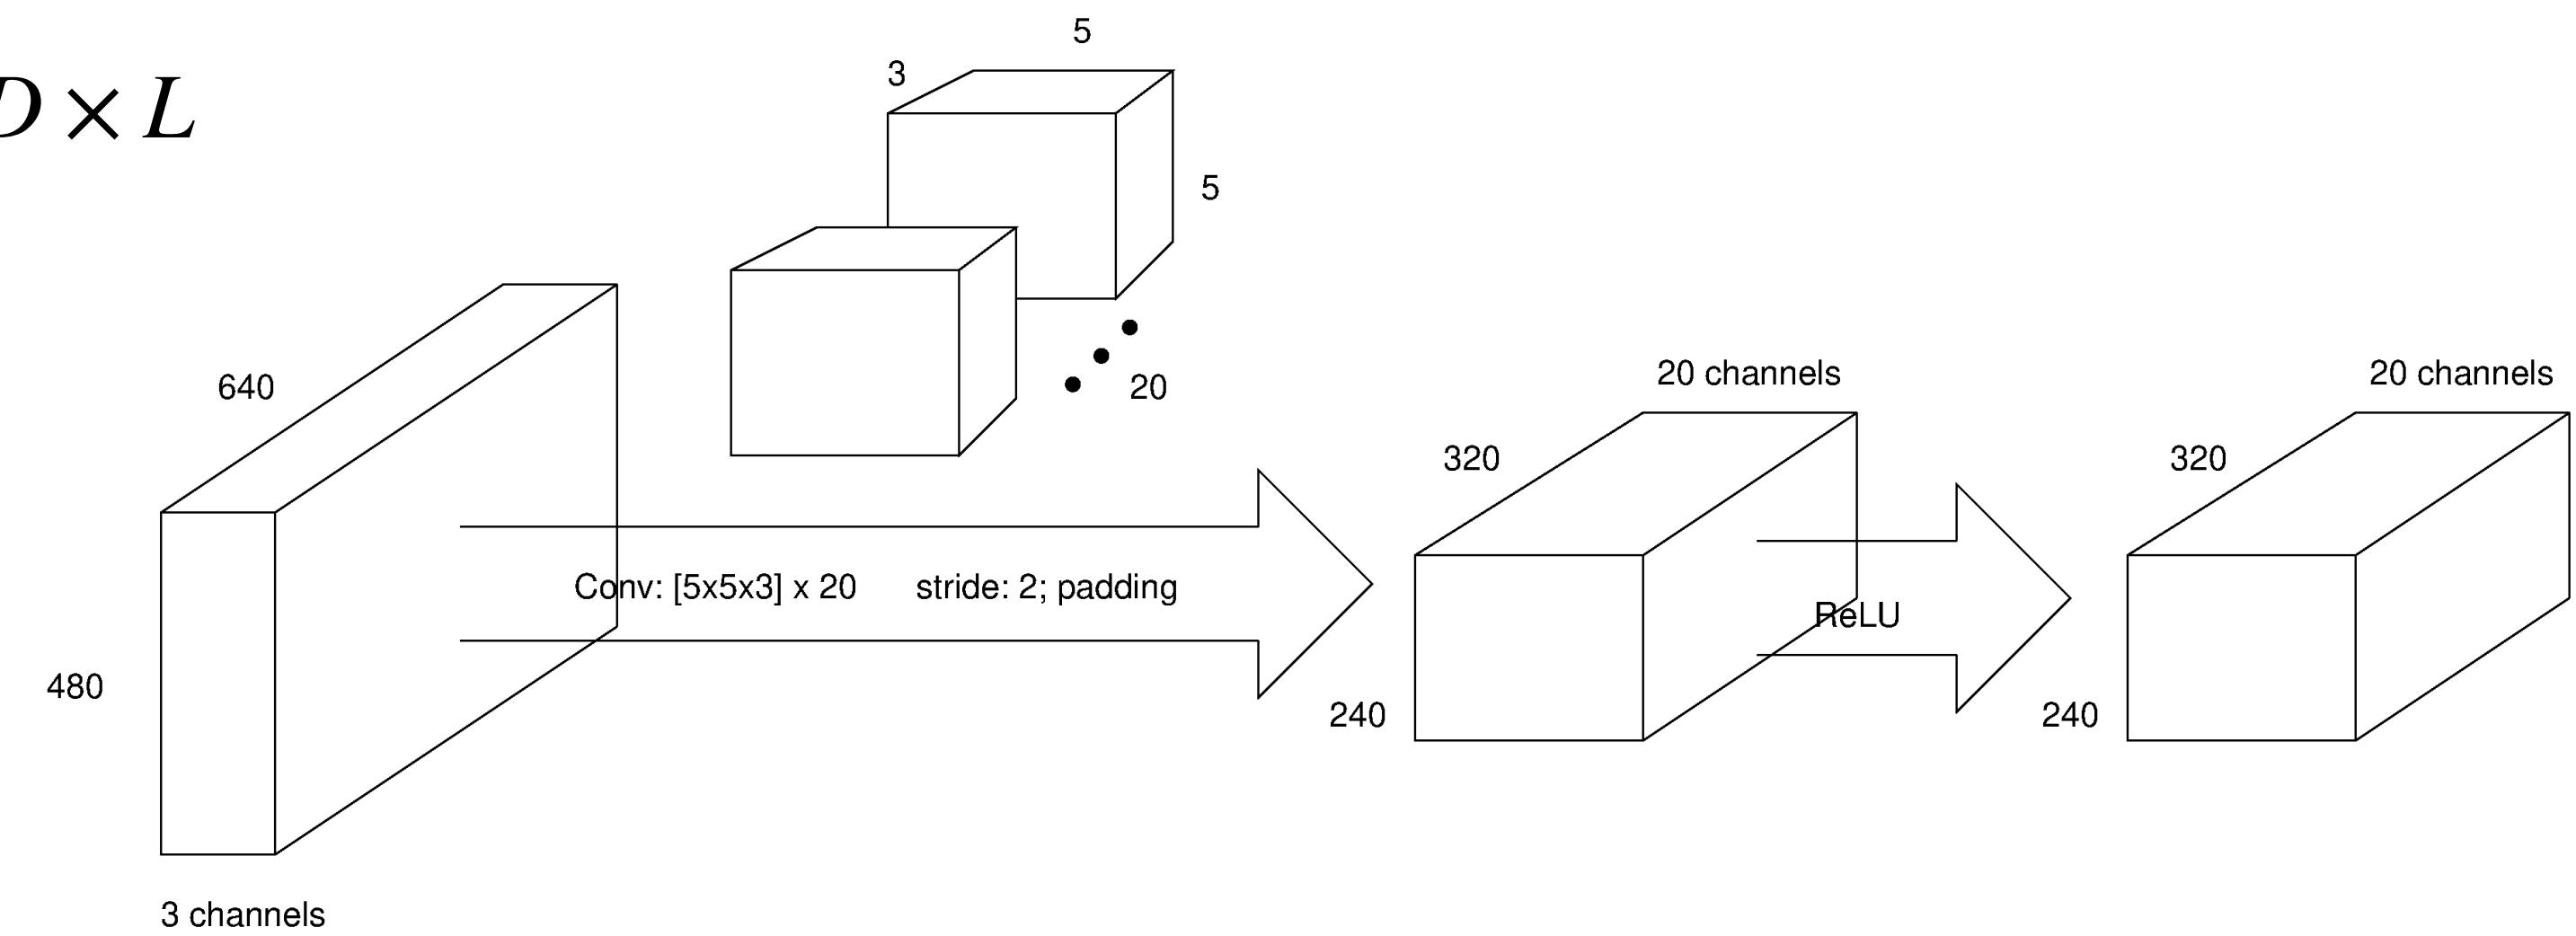
Convolutional Layer

- Linear Unit: $u = \sum_i w_i x_i + b$
- ReLU: $F(u) = \max(0, u)$
- Convolution: $N_{u,v} = \sum_{(i,j) \in K} I_{u+i, v+j} K_{i,j}$
- Linear unit implements convolution:
 - bias b : tune level of similarity for positive output
- Convolutional Layer
 - compute Feature Maps
 - learn kernel weights



Convolutional Layer

- Input: I of size $x \times y$ with D channels
 - 3-D: $x \times y \times D$
- Parameters: Shared among the $x \times y \times L$ units
 - Kernel K weights: dimensions: $k_i \times k_j \times D \times L$
 - Bias vector components: L
- Output: Stack of L convolutional responses
 - 3-dimensional: $x_n \times y_n \times L$
 - x_n, y_n : depend on strides and padding
 - ReLU layer



Pooling Layer

- Input: Feature map: $x_n \times y_n \times L$
- regions of adjacent pixels
 - same features activate for most
 - some features do not activate at all
- Pooling: capture patterns in groups of pixels
 - Patch of size $\{p_x \times p_y\}$ Stride lenght: s
 - Output: Feature map: $\lfloor \frac{x_n}{s} \rfloor \times \lfloor \frac{y_n}{s} \rfloor \times L$
 - Pooling function for inputs I that overlap patch
 - Maximum Pooling: maximum value in patch
 - Average Pooling: average value in patch
- Typical sequence of layers: Convolutional => ReLU => Pooling
- Pooling is robust

1	1	1	2
1	2	2	0
2	2	0	0
2	0	0	0

2	2
2	0

1.25	1.25
1.5	0

Max

Avg

Convolutional Layer

- Feature Maps
- From convolutional kernel to linear unit
- Convolutional Layer
- Pooling Layer

Applied Machine Learning

Convolutional Layer